

Client review with accountable AI

A working n8n Cloud prototype using 34 fictional accounts, connected data-table retrieval, OpenAI drafting and a human approval gate.

34

account results saved

24

drafted items; 10 need no outreach

1

current-version approval shown

What the system does

Element	Implemented behaviour
Trigger	A published two-minute schedule runs the demonstration without a fresh chat prompt. A batch holds up to eight accounts.
Inputs	Setup execution 93 seeded two Data Tables from embedded fictional CRM and delivery JSON. Scheduled processing retrieves and joins the records. No manual CSV import or live CRM connection was used.
Rules + AI	The latest configuration computes fixed risk conditions in code. The OpenAI node explains supplied decisions and drafts personalised emails or internal notes. Output validation remains in place.
Destination	Current results, source snapshots, versions and model identifiers are saved in a review queue. Separate audit records preserve processing and failure evidence.
Human gate	An authenticated form displays the saved version and source facts. A person judges accuracy, tone and suitability before approving that specific version. There is no sending node.

Saved results: all 34 accounts

Execution	Accounts	Version	Saved
98	C001-C008	1	8
234	C009-C016	2	8
270	C017-C024	3	8
272	C025-C032	1	8
276	C033-C034	2	2

Final saved classification: Green 10 | Amber 12 | Red 8 | Review Pending 4. All source snapshots match the fictional fixture. These results span configuration changes and recovery attempts; they are not one unchanged run. Recorded model: gpt-5.4-mini-2026-03-17.

A human decision, backed by a traceable version

The review form displayed **C002 / Beacon Health / version 1**, reference **GEN-98-C002-V1**, with the email and source facts. I recorded approval of this exact version through the form.

Approval recorded
Beacon Health — version 1
Ready for manual sending
Only this current version was approved.
No client email was sent.

Evidence: the displayed version reference and approval confirmation are preserved as paired screenshots. This document reports form-UI evidence; the raw approval-node audit JSON and execution ID were not collected.

A real failure informed the design

Observed in execution 238: the model labelled C017 Green despite a 14-day contact gap, which the fixed rubric defines as Amber. Validation rejected the batch. Audit row 37 recorded **RUN_FAILED** against the correct run; eight claimed rows became **Processing failed / Not ready**. A later recovery saved those accounts in execution 270.

Design response: FIX03 assigns explicit threshold calculations to code and gives the resulting policy decisions to the model for explanation and drafting. Earlier raw errors remain in the evidence. Agreement with supplied labels is **not independent model classification accuracy**.

What this demonstrates

Automatic processing beyond chat, persisted results for 34 synthetic accounts, a version-specific human approval, and one observed AI-validation failure that was logged and blocked. Drafts still require judgment: some unapproved emails use awkward phrasing or unnecessary internal thresholds.

Boundaries

No client messages were sent. Live CRM integration, a last-working-day calendar, business time savings, and return/revision/stale-form tests are outside the demonstrated scope. The packed workflow files are templates and patches, not a fresh export of every deployed setting. Successful outputs do not establish production readiness.

Project contribution

I configured and ran the n8n workflows, inspected saved outputs and failure evidence, and made the demonstrated approval decision. ChatGPT assisted with workflow design, synthetic cases, code, debugging and documentation.